ASSIGNMENT 02 REPORT

22I-1963 **Muhammad Abdurrehman** **DS-A**

Legal Clause Semantic Similarity using NLP

---

1. Introduction

This project addresses the challenge of identifying semantic similarity between legal clauses. Legal documents often express the same principle in different wordings, making it crucial to detect equivalence through semantic understanding rather than surface-level matching. The task focuses on developing models capable of measuring semantic relatedness between pairs of legal clauses using non-transformer deep learning architectures.

---

2. Dataset Description

The **Legal-Clause-Dataset** consists of multiple CSV files, each corresponding to a clause category (e.g., *acceleration.csv*, *access.csv*).
Each file contains two columns: clause_text and clause_type.

After preprocessing and cleaning:

- Total clause samples loaded: **[Add count after execution]**
- Clause pair dataset generated: **40,000 pairs** (balanced positive and negative pairs)
- Label distribution:
  - Similar (1): 19,946
  - Dissimilar (0): 20,054

Split	Samples	Percentage
Training	28,900	72.25%
Validation	5,100	12.75%

Test	6,000	15.00%

Vocabulary size after tokenization: **28,743**
Maximum sequence length: **256 tokens**

---

3. Data Preprocessing

- Text normalization: removal of extra whitespace, quotes, and numeric/time standardization (<NUM>, <TIME>).
- Lowercasing of all tokens for uniformity.
- Generation of positive pairs (same clause type) and negative pairs (different clause types).
- Tokenization using Keras Tokenizer, followed by padding to fixed length.

```
Pairs generated: 40000
Label distribution:
 label
0    20054
1    19946
Name: count, dtype: int64
```

[5]:

	text_a	text_b	label
0	entire agreement; amendment. this escrow agree...	default. each of the following events shall co...	0
1	termination without cause. company also may te...	compensation. for the services provided hereun...	0
2	restrictive covenants. (a) for a period of twe...	restrictive covenants. it cruise is not a part...	1
3	redemption. t+ <num> r <num> . <num> prior to ...	w i t n e s s e t h whereas the seller is the ...	0
4	terms. the terms of such extended loans shall ...	subsidiaries. <num> . <num> (a) ucc filing jur...	0
5	special terms and conditions of trust. tarantu...	binding effect. this agreement inures to the b...	0

---

4. Model Architectures

Model A: Siamese BiLSTM

- **Input:** Two sequences of length 256
- **Encoder:** Shared embedding (200-dim) + BiLSTM (128 units)
- **Combination:** Concatenation of element-wise difference and multiplication
- **Dense layers:** 128 → 64 → 1 (sigmoid)
- **Trainable parameters:** 6,225,217
- **Rationale:** Captures sequence-level semantic representations and computes relational similarity through feature interaction.

Model B: Attention-based BiLSTM

- **Input:** Two sequences of length 256

- **Encoder:** Shared embedding (200-dim) + BiLSTM (128 units, return sequences) + Custom Attention Layer (64 units)
- **Combination:** Concatenation of both embeddings and their absolute difference
- **Dense layers:** 128 → 64 → 1 (sigmoid)
- **Trainable parameters:** 6,208,962
- **Rationale:** Incorporates attention to focus on semantically salient tokens, improving interpretability and context weighting over long clauses.

---

5. Training Configuration

Parameter	Value
Optimizer	Adam
Loss	Binary Crossentropy
Metrics	Accuracy
Batch Size	128
Epochs	10
Early Stopping	Patience = 3
ReduceLROnPlateau	Factor = 0.5
Validation Split	15%

GPU acceleration (Kaggle environment) was utilized for efficient training.

---

6. Training Curves

(Insert screenshots of Cell 12 outputs here)
Figures show decreasing training/validation loss and increasing accuracy, indicating effective convergence and minimal overfitting.

---

7. Evaluation Metrics

Performance measured using:

- **Accuracy:** Percentage of correct predictions.
- **Precision:** Ratio of correctly predicted “similar” pairs to all predicted similar pairs.
- **Recall:** Fraction of actual similar pairs correctly identified.
- **F1-score:** Harmonic mean of precision and recall for balanced evaluation.
- **ROC-AUC:** Measures overall discriminative ability between similar/dissimilar pairs.

Rationale:
In legal NLP, **F1-score** and **ROC-AUC** are most appropriate since class balance and misclassification cost matter. High recall ensures relevant similar clauses are not missed, while high precision reduces false equivalence in legal analysis.

8. Quantitative Performance Comparison

(Insert actual test results table from Cell 14 here)

Model	Accuracy	Precision	Recall	F1-score	ROC-AUC	Training Time
Siamese BiLSTM	[Add]	[Add]	[Add]	[Add]	[Add]	[Add]
Attention BiLSTM	[Add]	[Add]	[Add]	[Add]	[Add]	[Add]

---

9. Qualitative Results

(Attach screenshots from Cell 15)
Below are representative examples of correctly and incorrectly matched pairs.

Correctly predicted similar clauses (sample):

1. “Restrictive covenants...” ↔ “Restrictive covenants...”
2. “Termination without cause...” ↔ “Termination clause...”

Incorrect predictions (sample):

1. Misclassified clauses with similar surface structure but different legal implications.
2. Failure to detect similarity when terminology differs but intent matches.

---

10. Discussion

- Both baselines show that traditional sequence models can capture semantic similarity to a reasonable degree without pretrained transformers.
- The **Attention BiLSTM** demonstrates improved interpretability and robustness for long clauses.
- For real-world deployment, **F1-score** and **ROC-AUC** should be prioritized since legal tasks value balanced detection of both positive and negative matches.
- Future work could include domain-adaptive pretraining or hybrid models incorporating syntactic dependencies.

11. Conclusion

This assignment successfully implemented two deep learning baselines for legal clause similarity detection without using pretrained transformer models. The architectures effectively learned semantic relationships within clauses, with attention mechanisms showing superior contextual understanding.
These findings highlight the importance of domain-specific modeling for legal text similarity and provide a foundation for future transformer-based or hybrid explorations.